

```
In [1]: import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Student Performance EDA

Introduction

- Dataset: Student performance data with ~395 records.
- Objective: Understand student demographics, family, study habits etc

```
In [3]: df=pd.read_csv('student_mat.csv')
```

```
In [4]: df.head()
```

Out[4]: **school;sex;age;address;famsize;Pstatus;Medu;Fedu;Mjob;Fjob;reason;guardian;traveltime;s**

0

1

2

3

4



```
In [5]: df=pd.read_csv('student_mat.csv', sep=';', quotechar='"')
```

```
In [6]: df.shape
```

Out[6]: (395, 33)

```
In [7]: df.head()
```

Out[7]:

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	fan
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	...	
1	GP	F	17	U	GT3	T	1	1	at_home	other	...	
2	GP	F	15	U	LE3	T	1	1	at_home	other	...	
3	GP	F	15	U	GT3	T	4	2	health	services	...	
4	GP	F	16	U	GT3	T	3	3	other	other	...	

5 rows × 33 columns



Data Overview

In [8]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
#   Column      Non-Null Count  Dtype
---  -
0   school      395 non-null    object
1   sex         395 non-null    object
2   age         395 non-null    int64
3   address     395 non-null    object
4   famsize     395 non-null    object
5   Pstatus     395 non-null    object
6   Medu        395 non-null    int64
7   Fedu        395 non-null    int64
8   Mjob        395 non-null    object
9   Fjob        395 non-null    object
10  reason      395 non-null    object
11  guardian    395 non-null    object
12  traveltime  395 non-null    int64
13  studytime   395 non-null    int64
14  failures    395 non-null    int64
15  schoolsup   395 non-null    object
16  famsup      395 non-null    object
17  paid        395 non-null    object
18  activities  395 non-null    object
19  nursery     395 non-null    object
20  higher      395 non-null    object
21  internet    395 non-null    object
22  romantic    395 non-null    object
23  famrel      395 non-null    int64
24  freetime    395 non-null    int64
25  goout       395 non-null    int64
26  Dalc        395 non-null    int64
27  Walc        395 non-null    int64
28  health      395 non-null    int64
29  absences    395 non-null    int64
30  G1          395 non-null    int64
31  G2          395 non-null    int64
32  G3          395 non-null    int64
dtypes: int64(16), object(17)
memory usage: 102.0+ KB

```

```
In [9]: df.describe()
```

Out[9]:

	age	Medu	Fedu	traveltime	studytime	failures	fam
count	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000
mean	16.696203	2.749367	2.521519	1.448101	2.035443	0.334177	3.944300
std	1.276043	1.094735	1.088201	0.697505	0.839240	0.743651	0.896000
min	15.000000	0.000000	0.000000	1.000000	1.000000	0.000000	1.000000
25%	16.000000	2.000000	2.000000	1.000000	1.000000	0.000000	4.000000
50%	17.000000	3.000000	2.000000	1.000000	2.000000	0.000000	4.000000
75%	18.000000	4.000000	3.000000	2.000000	2.000000	0.000000	5.000000
max	22.000000	4.000000	4.000000	4.000000	4.000000	3.000000	5.000000

In [10]: `df.isnull().sum()`

Out[10]:

school	0
sex	0
age	0
address	0
famsize	0
Pstatus	0
Medu	0
Fedu	0
Mjob	0
Fjob	0
reason	0
guardian	0
traveltime	0
studytime	0
failures	0
schoolsup	0
famsup	0
paid	0
activities	0
nursery	0
higher	0
internet	0
romantic	0
famrel	0
freetime	0
goout	0
Dalc	0
Walc	0
health	0
absences	0
G1	0
G2	0
G3	0

dtype: int64

In [11]:

```
cat_col=df.select_dtypes(include=['object']).columns.tolist()
num_col=df.select_dtypes(exclude=['object']).columns.tolist()
```

In [12]: `print(cat_col)`

```
['school', 'sex', 'address', 'famsize', 'Pstatus', 'Mjob', 'Fjob', 'reason', 'guardian', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery', 'higher', 'internet', 'romantic']
```

```
In [13]: print(num_col)
```

```
['age', 'Medu', 'Fedu', 'traveltime', 'studytime', 'failures', 'famrel', 'freetime', 'goout', 'Dalc', 'Walc', 'health', 'absences', 'G1', 'G2', 'G3']
```

```
In [14]: for col in cat_col[:5]:
          print(f"\nValuecounts for {col}:")
          print(df[col].value_counts())
```

Valuecounts for school:

school

GP 349

MS 46

Name: count, dtype: int64

Valuecounts for sex:

sex

F 208

M 187

Name: count, dtype: int64

Valuecounts for address:

address

U 307

R 88

Name: count, dtype: int64

Valuecounts for famsize:

famsize

GT3 281

LE3 114

Name: count, dtype: int64

Valuecounts for Pstatus:

Pstatus

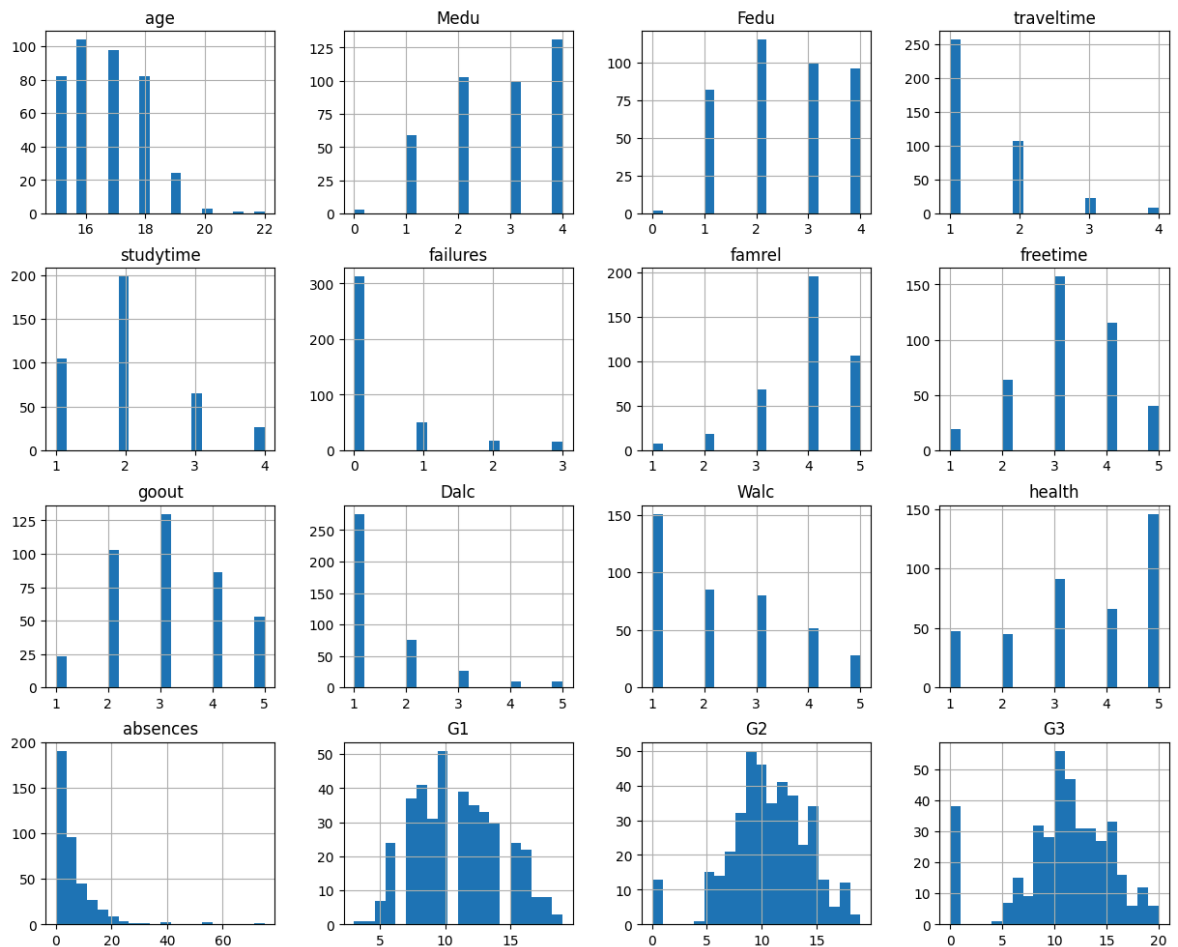
T 354

A 41

Name: count, dtype: int64

```
In [15]: df[num_col].hist(figsize=(15, 12), bins=20)
          plt.suptitle("Numerical Feature Distributions")
          plt.show()
```

Numerical Feature Distributions



- Total records: 395
- Columns: 33
- Data types: Mix of int and object
- Missing values: No missing values



Univariate Analysis

```
In [16]: #Numerical columns
print("🚀 Numerical Features Analysis\n")

for col in num_col:
    plt.figure(figsize=(12,5))

    #Histogram
    plt.subplot(1,2,1)
    sns.histplot(df[col],kde=True,bins=20,color='skyblue')
    plt.title(f"{col} - Distribution")


    #Boxplot
    plt.subplot(1,2,2)
    sns.boxplot(x=df[col],color='lightcoral')
    plt.title(f"{col} - Boxplot")

    plt.tight_layout()
```

```
plt.show()
```


```
print(f"{col} Summary Stats:\n", df[col].describe(), "\n")
```

📌 Numerical Features Analysis




age Summary Stats:

```
count    395.000000
mean     16.696203
std       1.276043
min       15.000000
25%      16.000000
50%      17.000000
75%      18.000000
max       22.000000
Name: age, dtype: float64
```



Medu Summary Stats:


```
count    395.000000
mean     2.749367
std       1.094735
min       0.000000
25%       2.000000
50%       3.000000
75%       4.000000
max       4.000000
Name: Medu, dtype: float64
```



Fedu Summary Stats:

```
count    395.000000
mean      2.521519
std       1.088201
min       0.000000
25%       2.000000
50%       2.000000
75%       3.000000
max       4.000000
```


Name: Fedu, dtype: float64



traveltime Summary Stats:

```
count    395.000000
mean      1.448101
std       0.697505
min       1.000000
25%       1.000000
50%       1.000000
75%       2.000000
max       4.000000
```


Name: traveltime, dtype: float64



studytime Summary Stats:

```
count    395.000000
mean      2.035443
std       0.839240
min       1.000000
25%       1.000000
50%       2.000000
75%       2.000000
max       4.000000
```


Name: studytime, dtype: float64



failures Summary Stats:

```
count    395.000000
mean      0.334177
std       0.743651
min       0.000000
25%       0.000000
50%       0.000000
75%       0.000000
max       3.000000
```


Name: failures, dtype: float64



famrel Summary Stats:

```
count    395.000000
mean      3.944304
std       0.896659
min       1.000000
25%      4.000000
50%      4.000000
75%      5.000000
max       5.000000
```


Name: famrel, dtype: float64



freetime Summary Stats:

```
count    395.000000
mean      3.235443
std       0.998862
min       1.000000
25%      3.000000
50%      3.000000
75%      4.000000
max       5.000000
```


Name: freetime, dtype: float64



goout Summary Stats:

```
count    395.000000
mean      3.108861
std       1.113278
min       1.000000
25%       2.000000
50%       3.000000
75%       4.000000
max       5.000000
```


Name: goout, dtype: float64



Dalc Summary Stats:

```
count    395.000000
mean      1.481013
std       0.890741
min       1.000000
25%       1.000000
50%       1.000000
75%       2.000000
max       5.000000
```


Name: Dalc, dtype: float64



Walc Summary Stats:

```
count    395.000000
mean      2.291139
std       1.287897
min       1.000000
25%       1.000000
50%       2.000000
75%       3.000000
max       5.000000
```


Name: Walc, dtype: float64



health Summary Stats:

```
count    395.000000
mean      3.554430
std       1.390303
min       1.000000
25%       3.000000
50%       4.000000
75%       5.000000
max       5.000000
```

Name: health, dtype: float64




absences Summary Stats:

```

count    395.000000
mean      5.708861
std       8.003096
min       0.000000
25%      0.000000
50%      4.000000
75%      8.000000
max      75.000000
Name: absences, dtype: float64

```




G1 Summary Stats:

```


count    395.000000
mean     10.908861
std      3.319195
min      3.000000
25%      8.000000
50%     11.000000
75%     13.000000
max     19.000000
Name: G1, dtype: float64

```



G2 Summary Stats:

```
count    395.000000
mean      10.713924
std        3.761505
min         0.000000
25%         9.000000
50%        11.000000
75%        13.000000
max        19.000000
Name: G2, dtype: float64
```



G3 Summary Stats:

```
count    395.000000
mean      10.415190
std        4.581443
min         0.000000
25%         8.000000
50%        11.000000
75%        14.000000
max        20.000000
Name: G3, dtype: float64
```

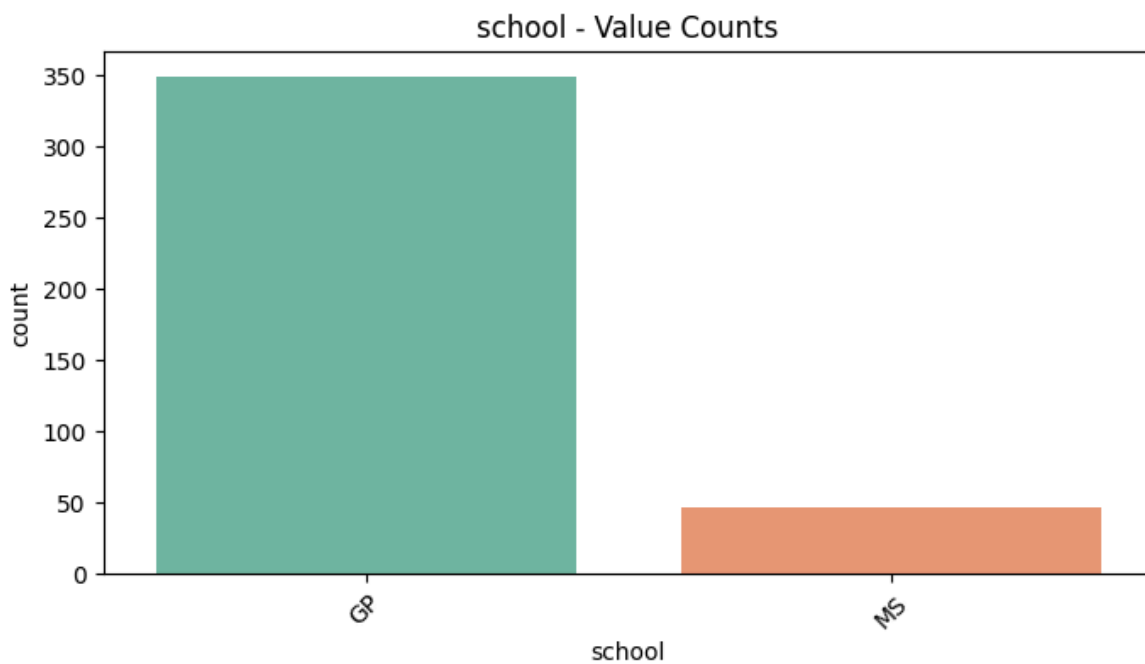
```
In [17]: # 2. Categorical Features
print("✦ Categorical Features Analysis\n")

for col in cat_col:
    plt.figure(figsize=(8,4))
    sns.countplot(x=df[col], order=df[col].value_counts().index, palette="Set2")
    plt.title(f"{col} - Value Counts")
    plt.xticks(rotation=45)
```

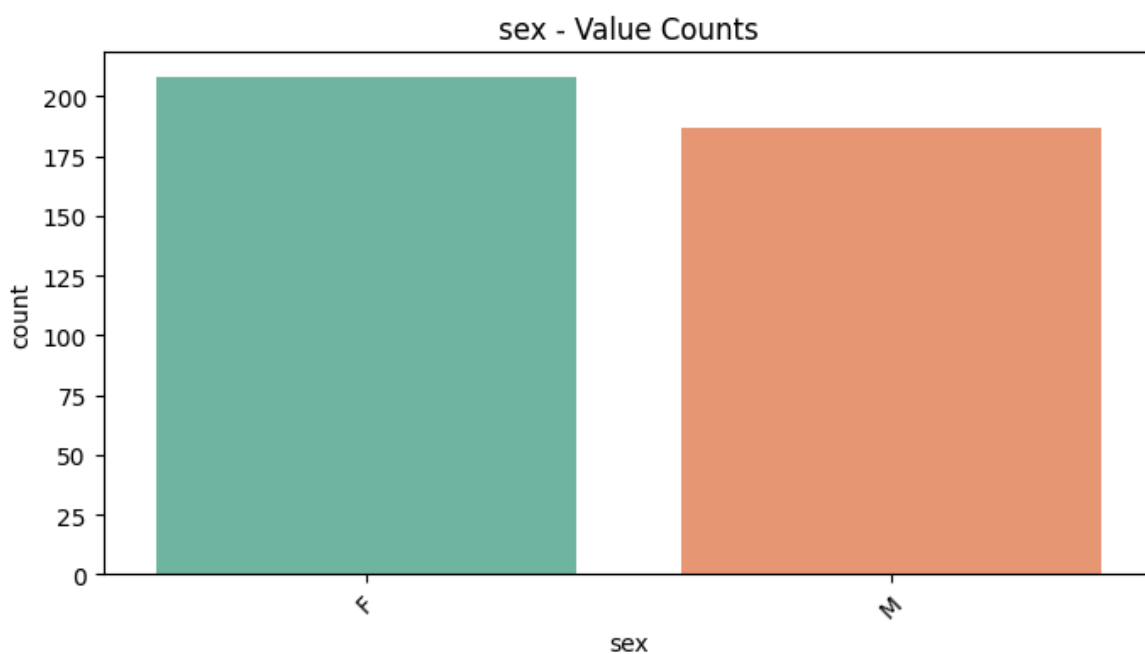
```
plt.show()
```

```
print(f"Value Counts for {col}:\n", df[col].value_counts(), "\n")
```

📌 Categorical Features Analysis



```
Value Counts for school:  
school  
GP    349  
MS     46  
Name: count, dtype: int64
```



```
Value Counts for sex:  
sex  
F    208  
M    187  
Name: count, dtype: int64
```




Value Counts for address:

address

U 307

R 88

Name: count, dtype: int64




Value Counts for famsize:

famsize

GT3 281


LE3 114

Name: count, dtype: int64



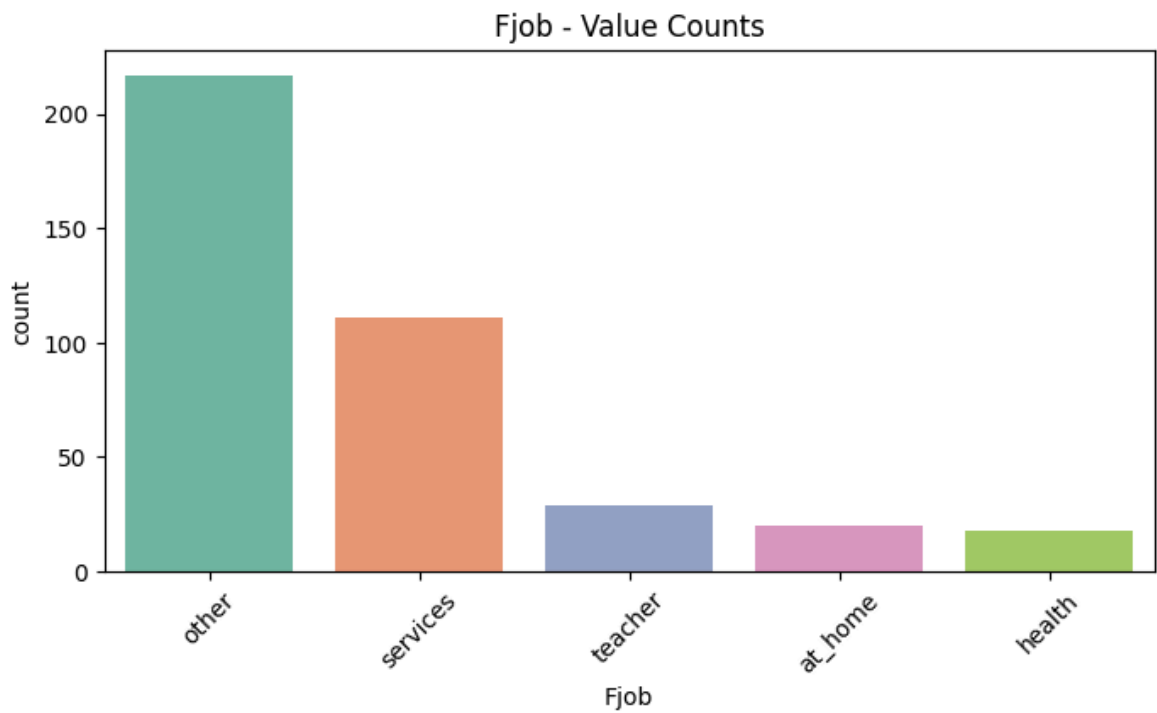
Value Counts for Pstatus:

```
Pstatus
T      354
A       41
Name: count, dtype: int64
```



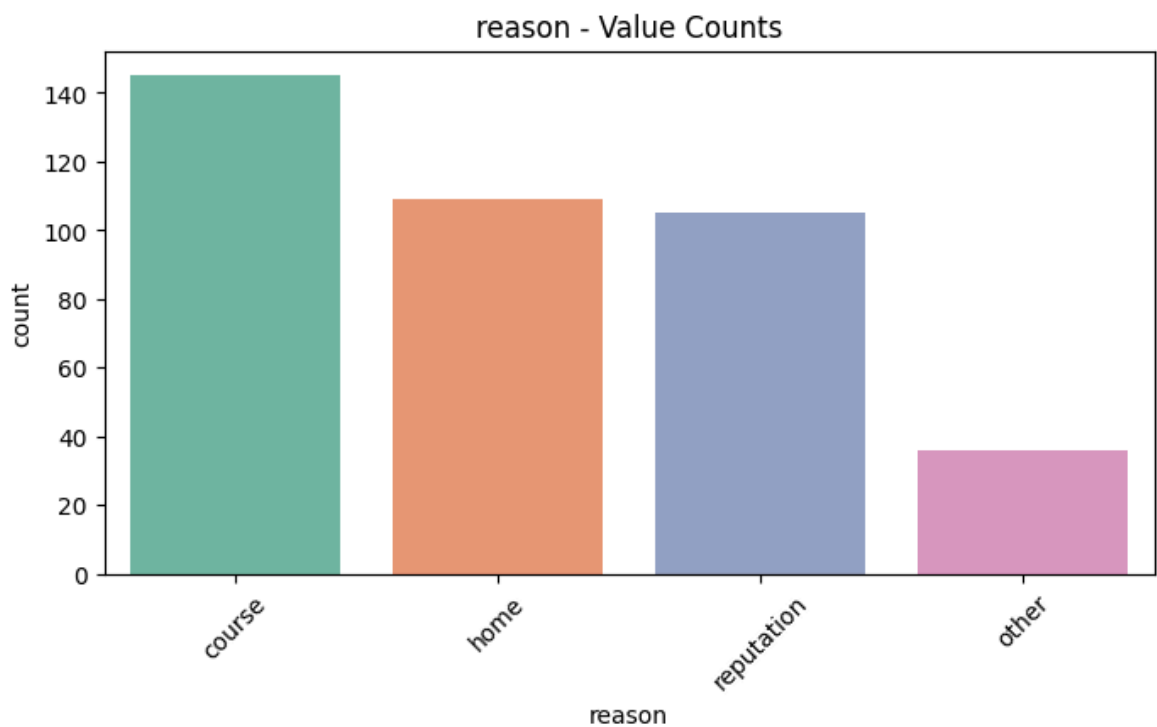
Value Counts for Mjob:

```
Mjob
other      141
services   103
at_home     59
teacher     58
health      34
Name: count, dtype: int64
```



Value Counts for Fjob:

```
Fjob
other      217
services   111
teacher     29
at_home    20
health     18
Name: count, dtype: int64
```



Value Counts for reason:

reason

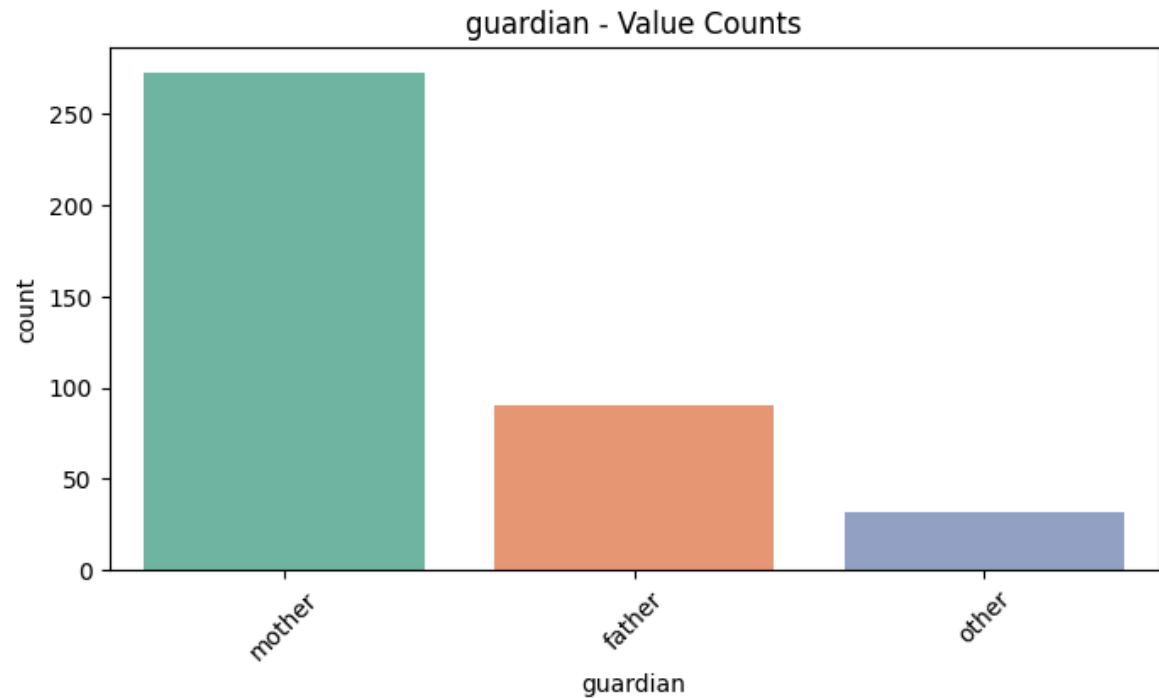
course 145

home 109

reputation 105

other 36

Name: count, dtype: int64



Value Counts for guardian:

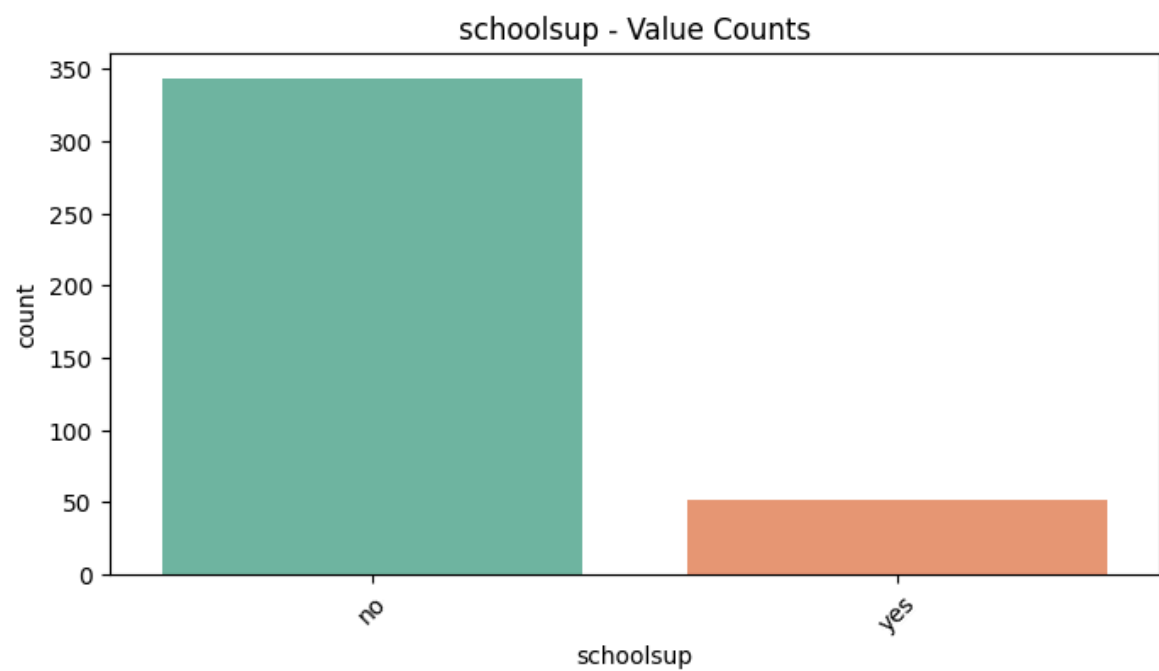
guardian

mother 273

father 90

other 32

Name: count, dtype: int64



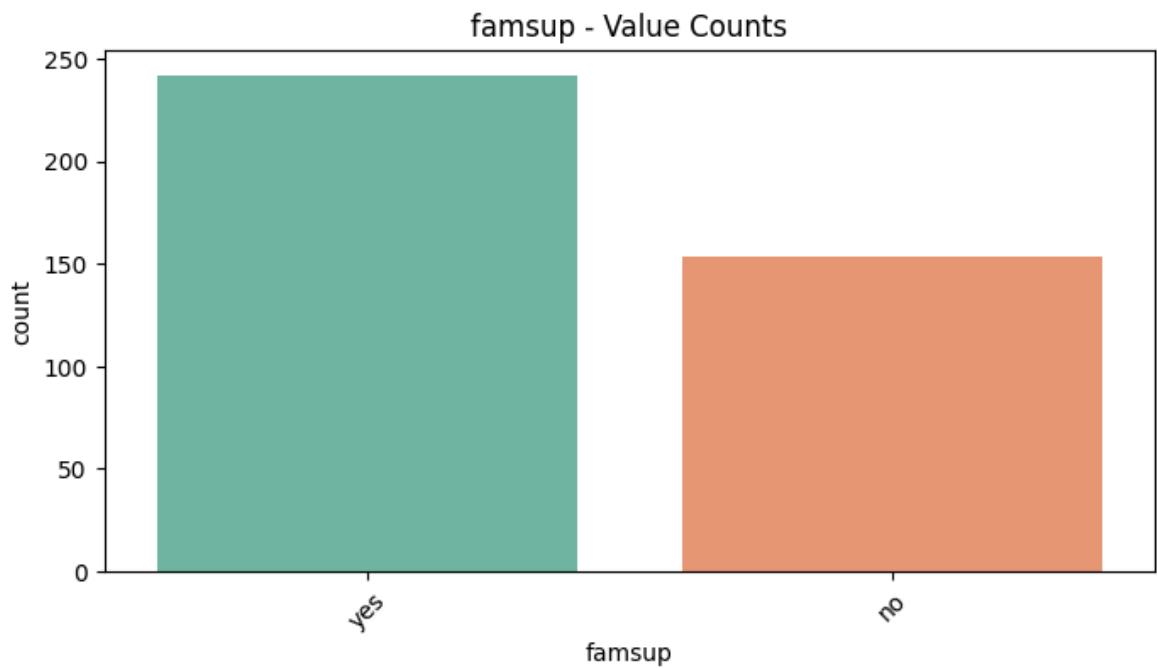
Value Counts for schoolsup:

schoolsup

no 344

yes 51

Name: count, dtype: int64



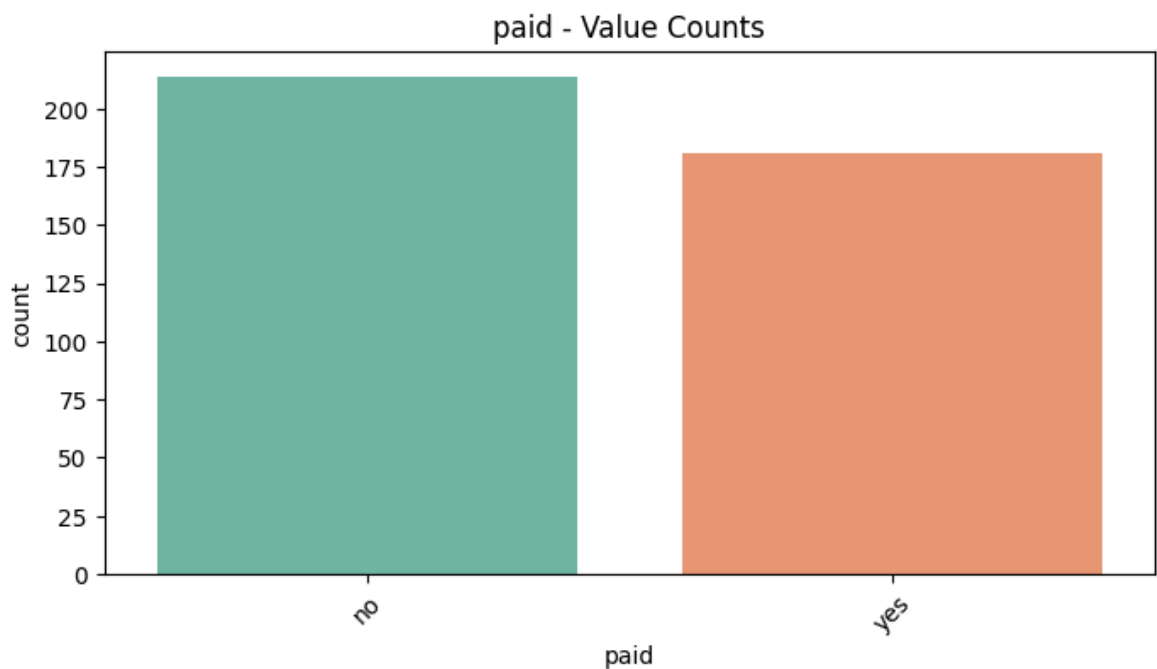
Value Counts for famsup:

famsup

yes 242

no 153

Name: count, dtype: int64



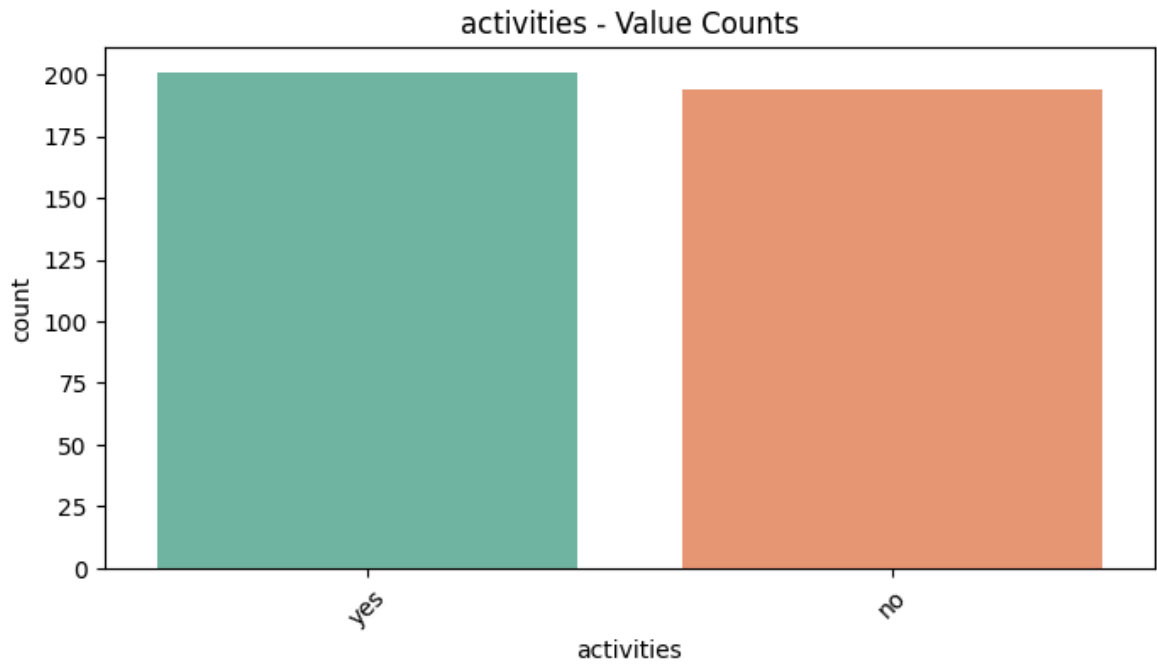
Value Counts for paid:

paid

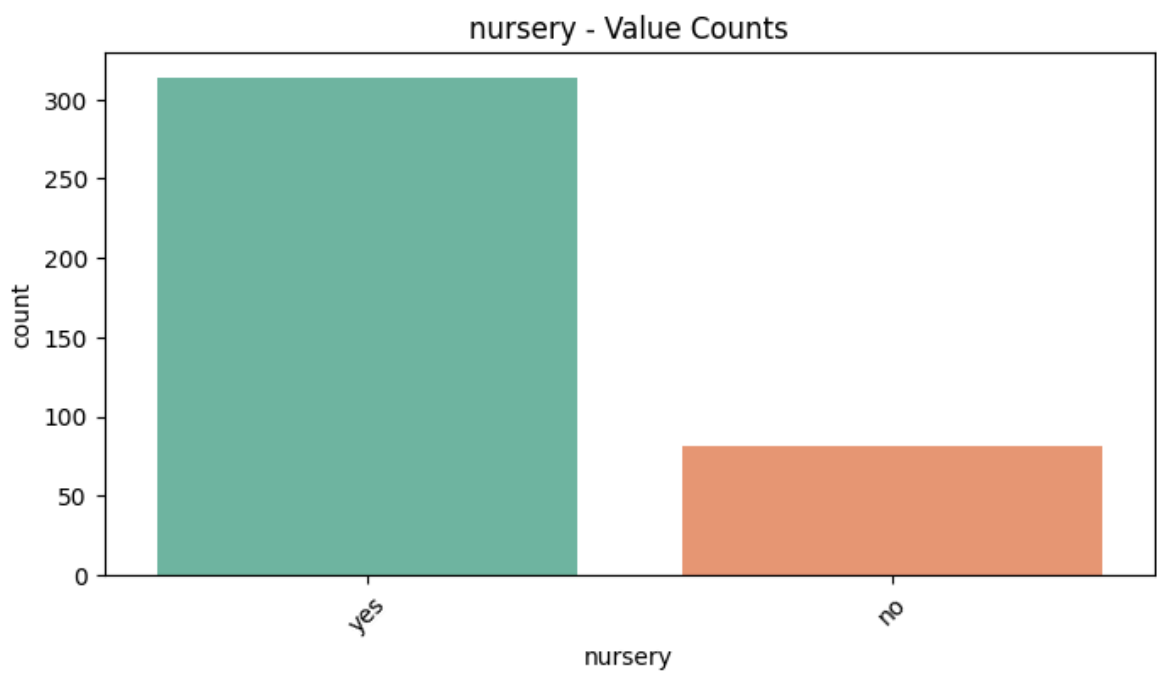
no 214

yes 181

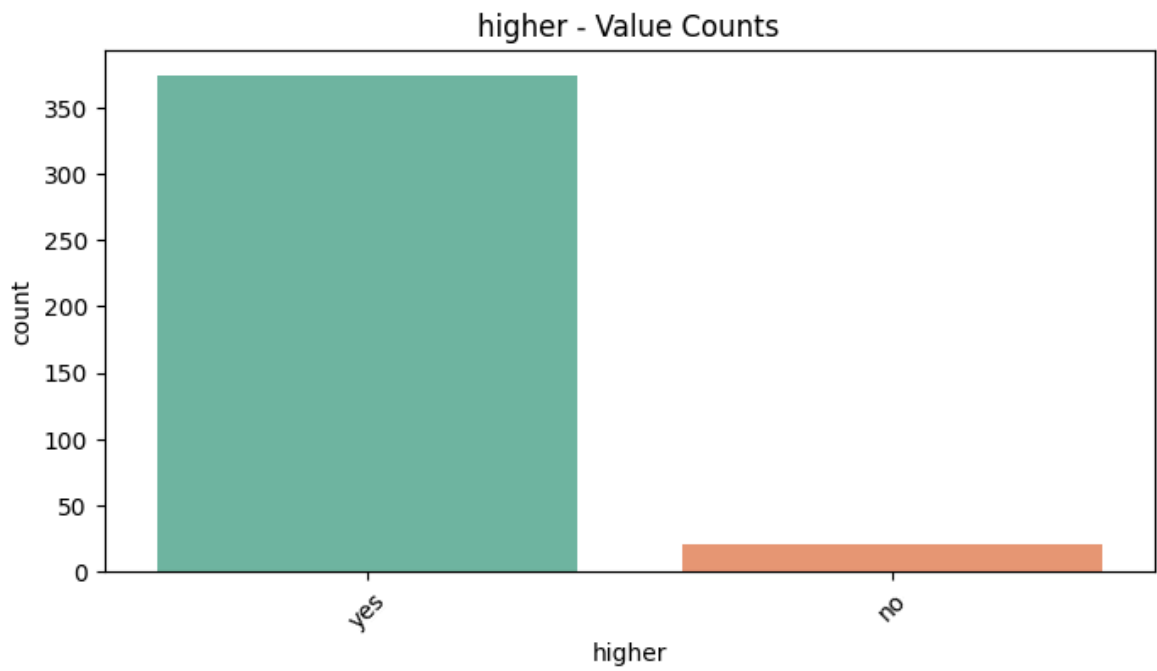
Name: count, dtype: int64



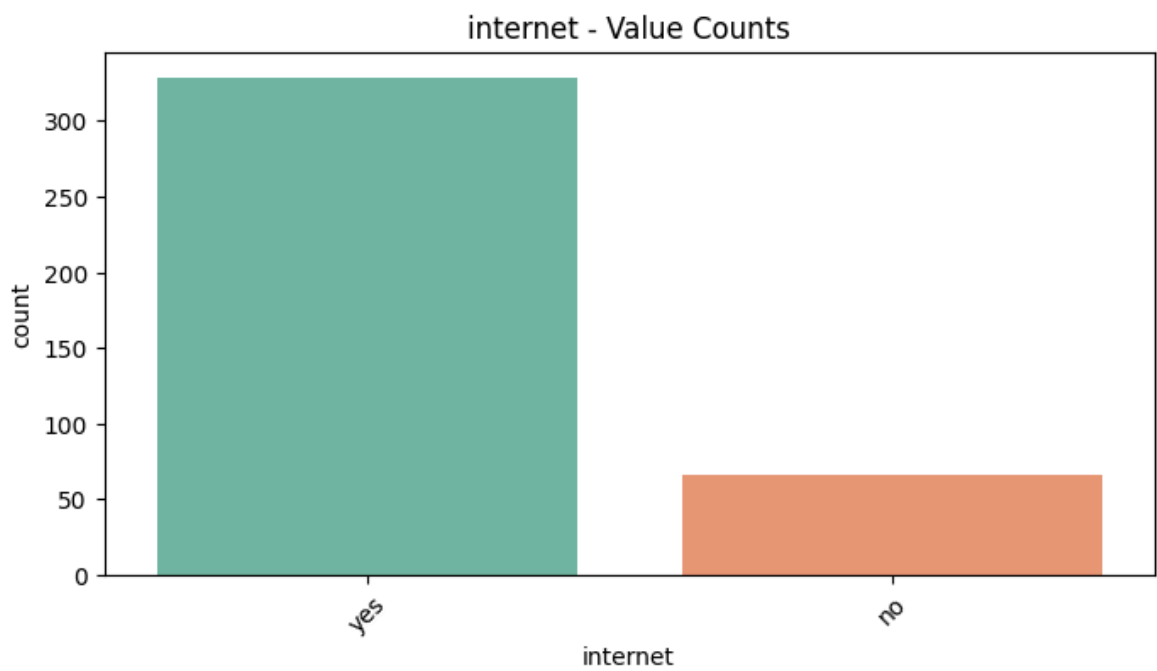
Value Counts for activities:
activities
yes 201
no 194
Name: count, dtype: int64




Value Counts for nursery:
nursery
yes 314
no 81
Name: count, dtype: int64



Value Counts for higher:
higher
yes 375
no 20
Name: count, dtype: int64



Value Counts for internet:
internet
yes 329
no 66
Name: count, dtype: int64



Value Counts for romantic:
romantic
no 263
yes 132
Name: count, dtype: int64

✓ Observations (Univariate EDA)

Categorical

- School → Majority students belong to GP, fewer from MS.
- Sex → Fairly balanced, but females slightly more than males.
- Address → More students live in Urban (U) areas compared to Rural (R).
- Family Size → Larger families (GT3) more common than smaller (LE3).
- Parent Status (Pstatus) → More students' parents are together (T) than apart (A).
- Guardian → Mother is most common guardian, followed by father, then other.
- Activities, Nursery, Internet →
- Participation in activities: almost equal yes/no
- Nursery attendance: more had nursery schooling
- Internet access: more have it than not

Numerical


- Age → Concentrated around mid-teen range (15–18).

- Absences → Heavily skewed with a few extreme outliers (some students missing dozens of classes).
- Grades (G1, G2, G3) → Not normally distributed, tend to cluster toward the lower end (many students performing below average).
- Other numeric features (studytime, freetime, goout, Dalc, Walc, health) are ordinal scales (discrete 1–5 ratings) rather than continuous numbers — so treat them almost like categorical when analyzing.



Bivariate Analysis

```
In [18]: target = "G3"
# 1. Correlation Heatmap (numerical vs numerical, including target)
plt.figure(figsize=(12,8))
sns.heatmap(df[num_col].corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap (Numerical Features)")
plt.show()
```



```
In [19]: # 2. Numerical Features vs Target (scatter + boxplot)
for col in [c for c in num_col if c != target]:
    plt.figure(figsize=(12,5))







    # Scatterplot
    plt.subplot(1,2,1)
    sns.scatterplot(x=df[col], y=df[target], alpha=0.6)
    plt.title(f"{col} vs {target} (Scatter)")

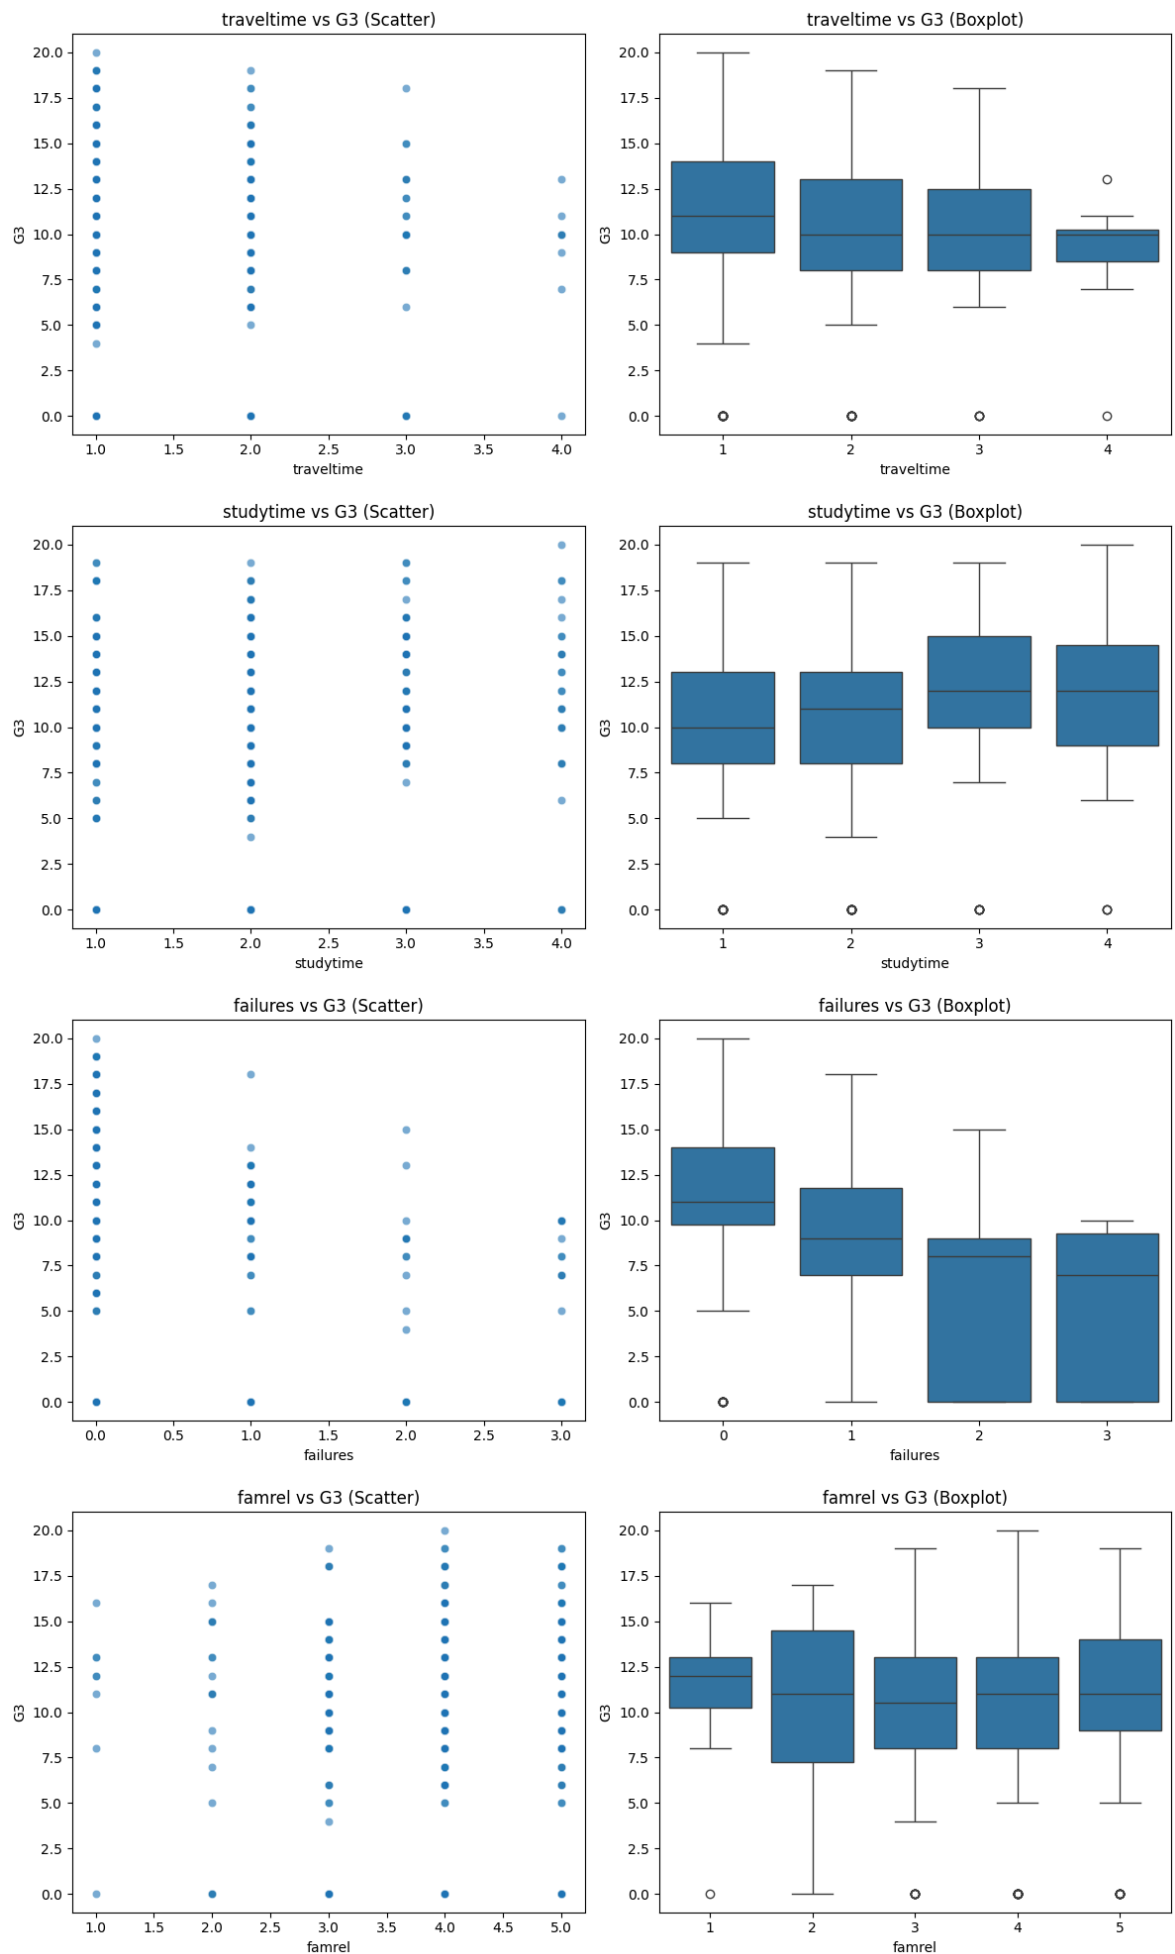
    # Boxplot (discretize col if it's ordinal like studytime, freetime, etc.)
```

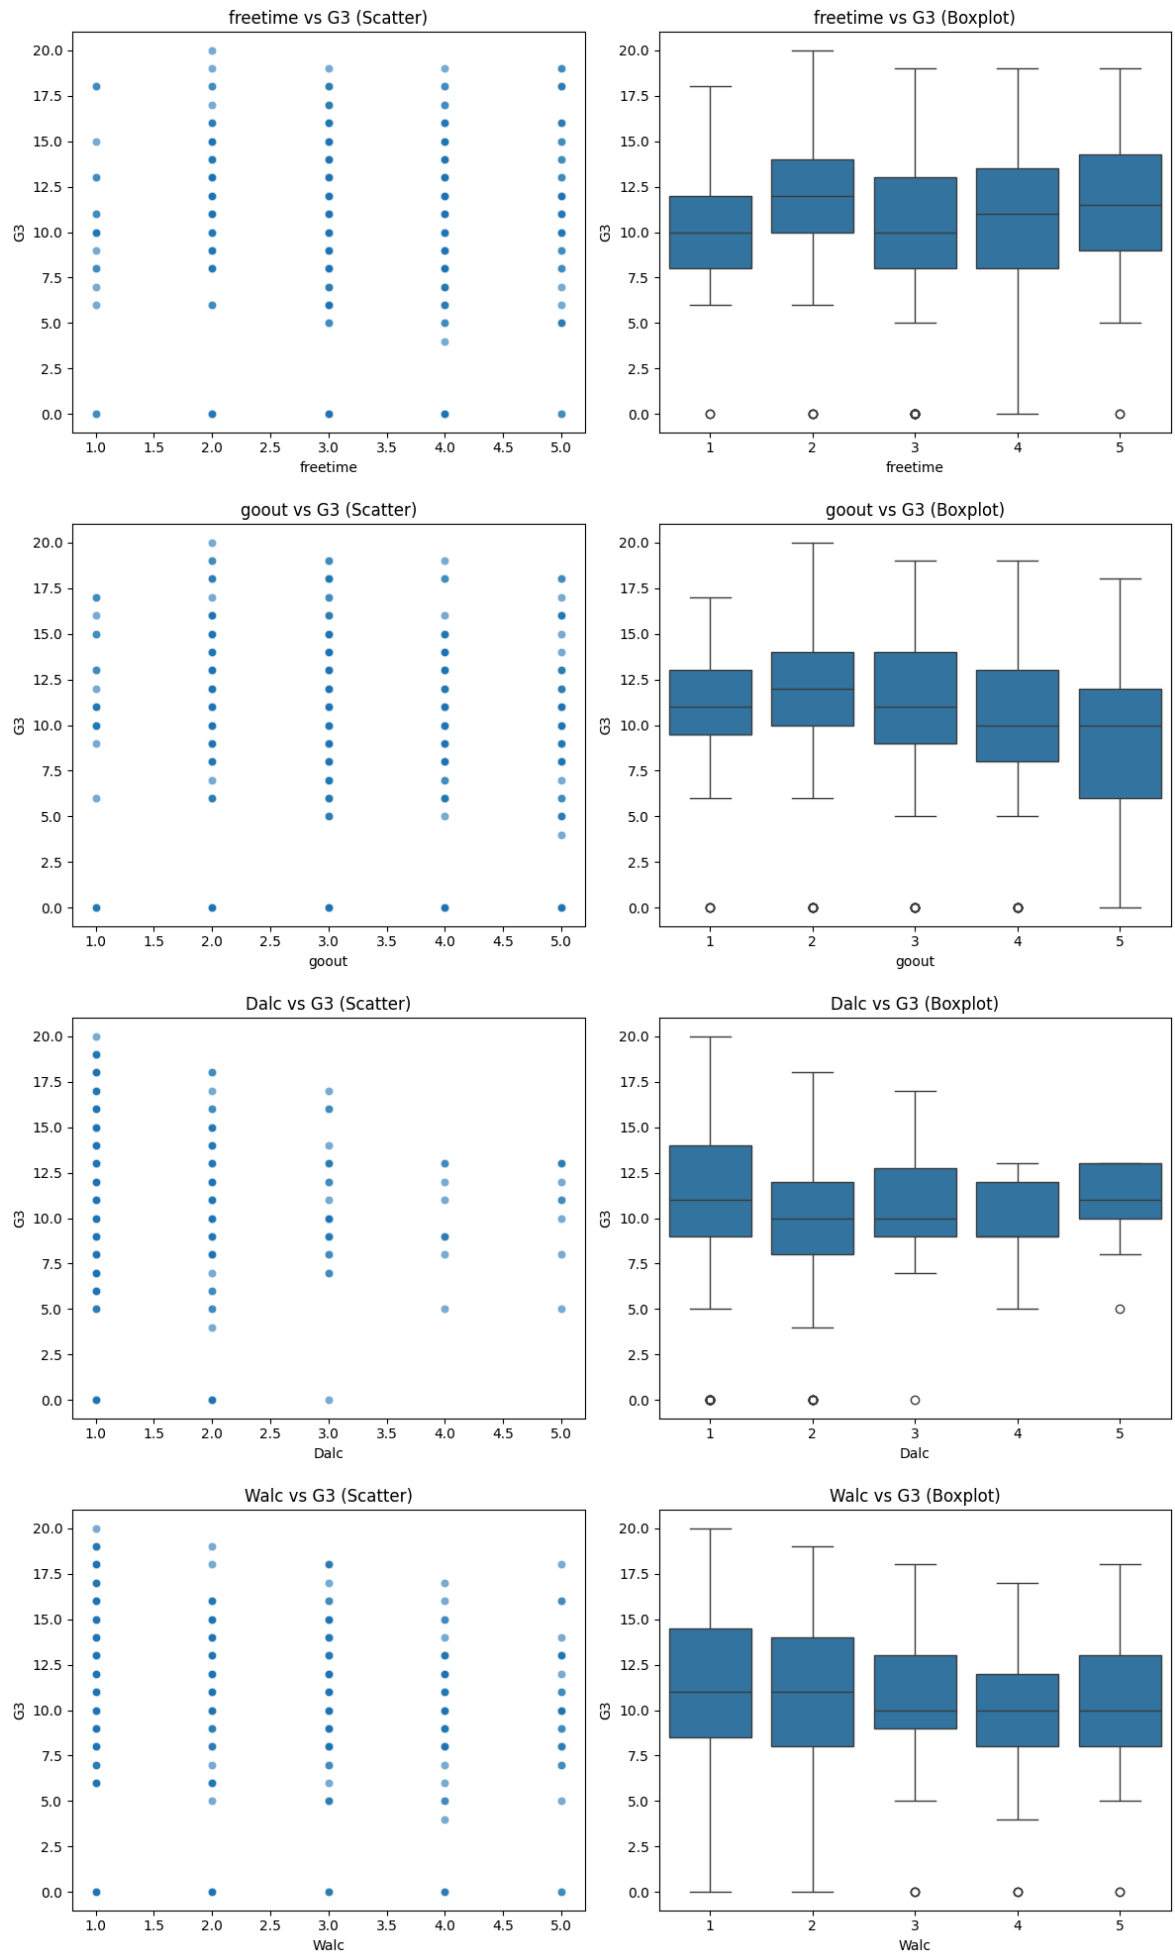


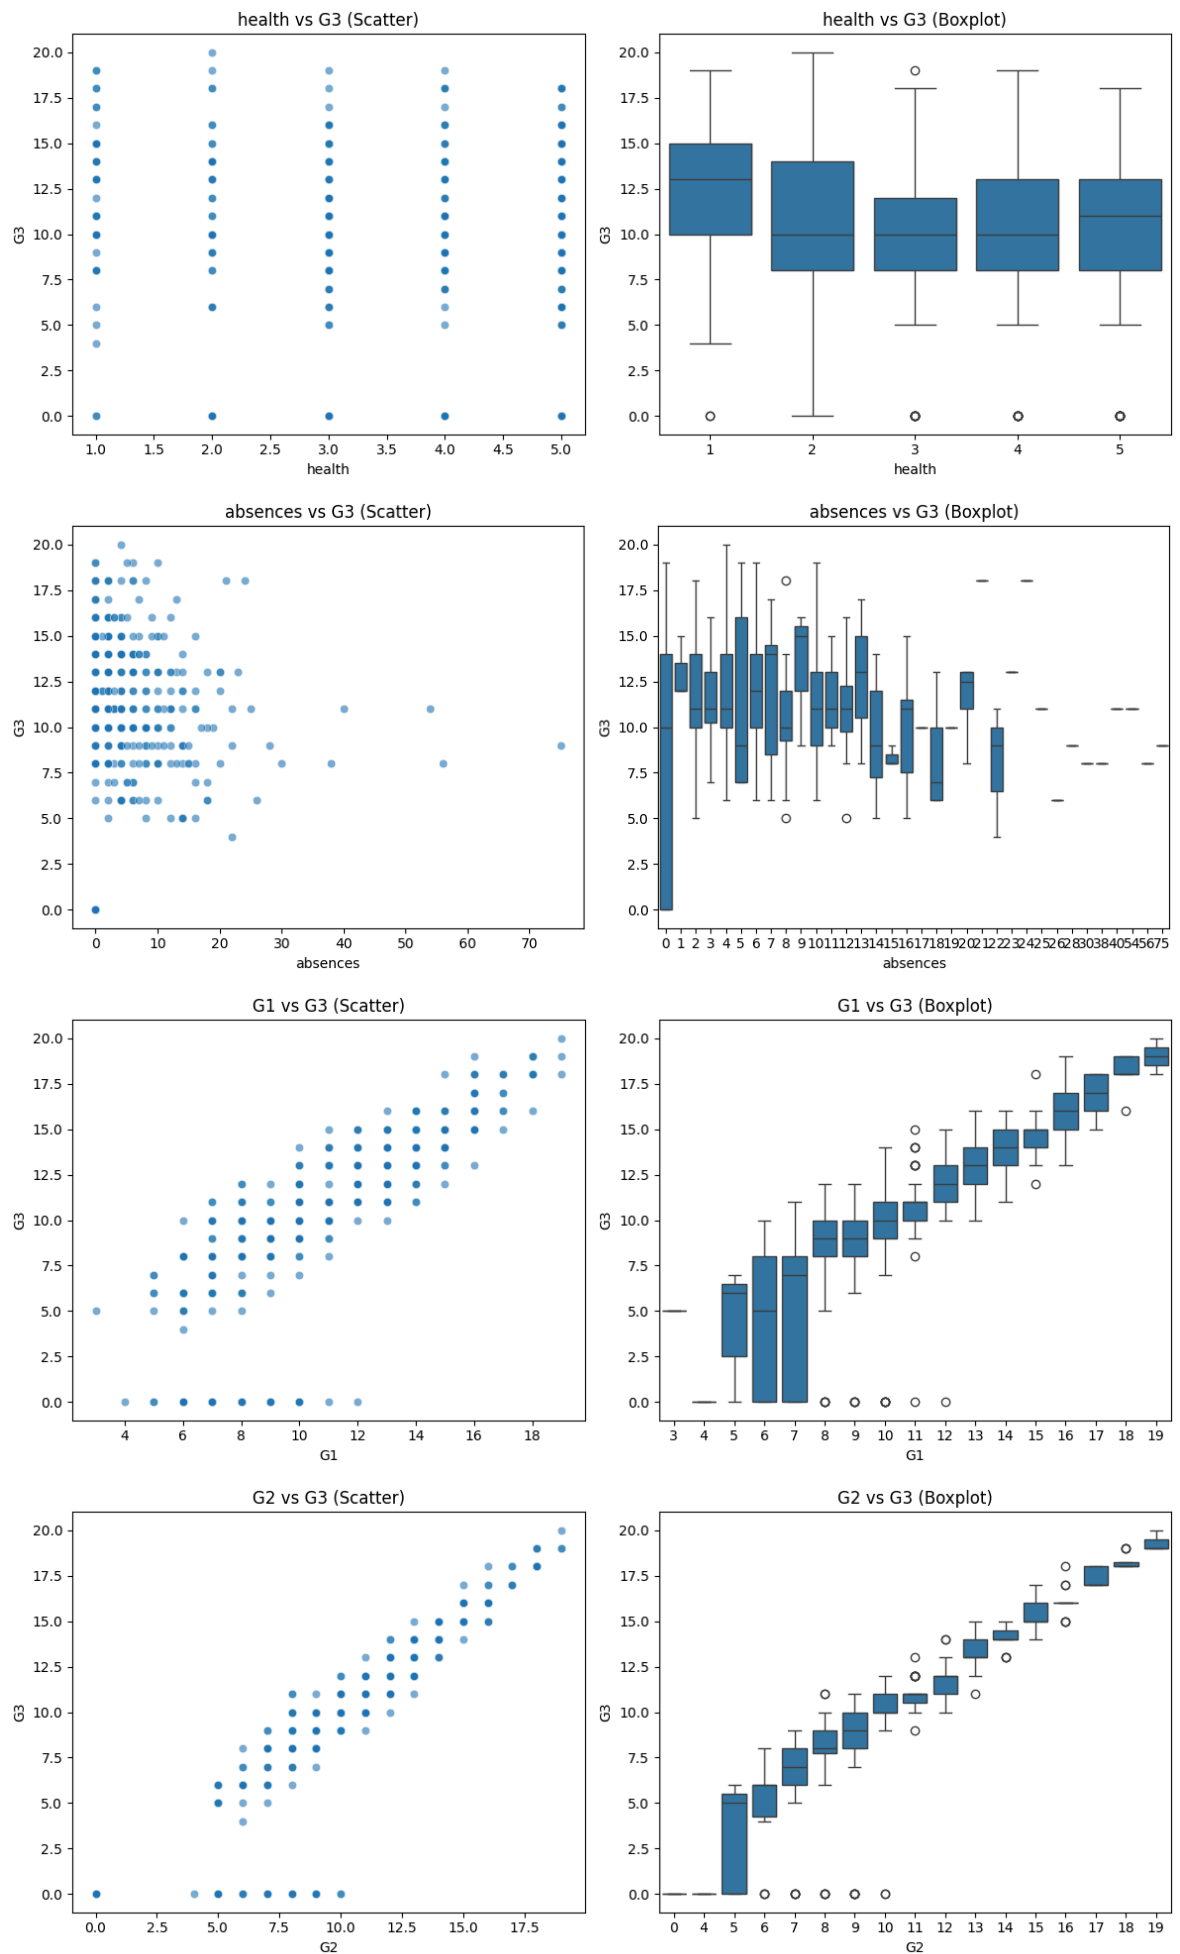
```
plt.subplot(1,2,2)
sns.boxplot(x=df[col], y=df[target])
plt.title(f"{col} vs {target} (Boxplot)")

plt.tight_layout()
plt.show()
```





















```
In [20]: for col in cat_col:
plt.figure(figsize=(10,5))
sns.boxplot(x=df[col], y=df[target], palette="Set2")
plt.title(f"{col} vs {target}")
plt.xticks(rotation=45)
plt.show()
```



















Bivariate Analysis – Notes

◆ Correlation Heatmap

G1 & G2 are very strongly correlated with G3 (final grade) → as expected, past performance is the best predictor of future performance.

Other numerical variables (studytime, freetime, health, etc.) show little to no linear correlation with G3.

◆ Numerical vs Target (G3)

Studytime vs G3 → Higher studytime categories are associated with slightly higher grades, but not a strong trend.

Freetime vs G3 → No clear relationship; grades are spread across all freetime levels.

Goout, Dalc (workday alcohol), Walc (weekend alcohol) → Higher values generally linked with slightly lower grades, but effect size is small.

Absences vs G3 → Students with very high absences have lower grades, though many with low absences also perform poorly → not a clean predictor.

◆ Categorical vs Target (G3)

School (GP vs MS) → Students from GP school slightly outperform those from MS on average.

Sex → Small difference; females tend to have slightly higher grades, but overlap is large.

Address (Urban vs Rural) → Urban students show a slight advantage, but not dramatic.

Guardian → Little difference across categories (mother, father, other).

Parental education & jobs → Some positive trend (higher education = slightly higher grades), but not very strong.

Activities, Internet, Nursery → No significant difference in G3 observed.

✓ Key Takeaways (Bivariate)

Strong Predictors: G1, G2 (past grades).

Moderate Predictors: Absences (negative effect), studytime (slightly positive), alcohol consumption (slightly negative).

Weak/No Clear Relationship: freetime, health, guardian, activities, internet access.

In []: